SimpleS3

by Nicholas Tierney

Abstract Writing functions in R is an important skill for an R programmer. Many R programmers are adept at creating their own functions, but do not use R's S3 methods. S3 methods allow for functions to be generalised across different classes and are easy to implement, once you know how. This guide is simple and targeted, serving to explain S3 methods so that users can create their own.

Note

This file is only a basic article template. For full details of *The R Journal* style and information on how to prepare your article for submission, see the Instructions for Authors.

Introduction

A standard principle of programming is DRY - Don't Repeat Yourself [ref]. Under this axiom, the copying and pasting of the same or similar code (copypasta), is avoided and replaced with a function. Having one function to replace several of the same or similar coded sections simplifies code maintenance as it means that only one section of code needs to be maintained, instead of several. This means that if the code breaks, then one simply needs to update the function, rather than finding all of the coded sections that are now broken.

S3 methods in the R programming language are a way of writing functions in R that do different things for objects of different classes. S3 methods are so named as the methods shipped with the release of the third version of the "S" programming language, which R was heavily based upon [reference]. Hence, methods for S 3.0 = S3 Methods.

The function `summary()` is an S3 method. When applied to an object of class "dataframe", `summary` shows descriptive statistics (Mean, SD, etc.) for each variable:

```
summary(iris)
```

```
#>   Sepal.Length    Sepal.Width     Petal.Length    Petal.Width
#>  Min.   :4.300   Min.   :2.000   Min.   :1.000   Min.   :0.100
#>  1st Qu.:5.100   1st Qu.:2.800   1st Qu.:1.600   1st Qu.:0.300
#>  Median :5.800   Median :3.000   Median :4.350   Median :1.300
#>  Mean   :5.843   Mean   :3.057   Mean   :3.758   Mean   :1.199
#>  3rd Qu.:6.400   3rd Qu.:3.300   3rd Qu.:5.100   3rd Qu.:1.800
#>  Max.   :7.900   Max.   :4.400   Max.   :6.900   Max.   :2.500
#>        Species
#>  setosa    :50
#>  versicolor:50
#>  virginica :50
#>
#>
#>
```

summary also performs differently when applied to different object, take a linear model, for example:

```
lm_iris <- lm(Sepal.Length ~ Sepal.Width, data = iris)
```

```
summary(lm_iris)
```

```
#>
#> Call:
#> lm(formula = Sepal.Length ~ Sepal.Width, data = iris)
#>
#> Residuals:
#>     Min      1Q  Median      3Q     Max
#> -1.5561 -0.6333 -0.1120  0.5579  2.2226
#>
#> Coefficients:
#>             Estimate Std. Error t value Pr(>|t|)
```

```
#> (Intercept)   6.5262     0.4789   13.63    <2e-16 ***
#> Sepal.Width  -0.2234     0.1551   -1.44     0.152
#> ---
#> Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
#>
#> Residual standard error: 0.8251 on 148 degrees of freedom
#> Multiple R-squared:  0.01382,    Adjusted R-squared:  0.007159
#> F-statistic: 2.074 on 1 and 148 DF,  p-value: 0.1519
```

summary produces a description of the linear model, describing how it was called (call), as well as the residuals, coefficients, t-values, p-values, R^2, and more. This output is **completely** different to the information output from summary used for the iris dataframe.

So how does the same function, summary perform differently for different objects? The answer is that R is sneaky, and *hides* information. There are in fact, many different summary functions. For example:

- summary.lm
- summary.data.frame
- summary.Date
- summary.matrix

Being an S3 method, summary calls the appropriate function based upon the class of the object it operates on. So using summary on an object of class "Date" will evoke the function, summary.Date. **But all you need to do is type summary**, and the S3 method does the rest.

To further illustrate, using summary on the iris data will actually call the function summary.data.frame, since iris is of class data.frame. We can find the class of an object using class

```
class(iris)

#> [1] "data.frame"

summary.data.frame(iris)

#>   Sepal.Length    Sepal.Width     Petal.Length    Petal.Width
#>   Min.   :4.300   Min.   :2.000   Min.   :1.000   Min.   :0.100
#>   1st Qu.:5.100   1st Qu.:2.800   1st Qu.:1.600   1st Qu.:0.300
#>   Median :5.800   Median :3.000   Median :4.350   Median :1.300
#>   Mean   :5.843   Mean   :3.057   Mean   :3.758   Mean   :1.199
#>   3rd Qu.:6.400   3rd Qu.:3.300   3rd Qu.:5.100   3rd Qu.:1.800
#>   Max.   :7.900   Max.   :4.400   Max.   :6.900   Max.   :2.500
#>        Species
#>   setosa    :50
#>   versicolor:50
#>   virginica :50
#>
#>
#>
```

which is the same as

```
summary(iris)

#>   Sepal.Length    Sepal.Width     Petal.Length    Petal.Width
#>   Min.   :4.300   Min.   :2.000   Min.   :1.000   Min.   :0.100
#>   1st Qu.:5.100   1st Qu.:2.800   1st Qu.:1.600   1st Qu.:0.300
#>   Median :5.800   Median :3.000   Median :4.350   Median :1.300
#>   Mean   :5.843   Mean   :3.057   Mean   :3.758   Mean   :1.199
#>   3rd Qu.:6.400   3rd Qu.:3.300   3rd Qu.:5.100   3rd Qu.:1.800
#>   Max.   :7.900   Max.   :4.400   Max.   :6.900   Max.   :2.500
#>        Species
#>   setosa    :50
#>   versicolor:50
#>   virginica :50
#>
#>
#>
```

And using summary on the linear model object, `lm_iris` performs:

```
summary.lm(lm_iris)

#>
#> Call:
#> lm(formula = Sepal.Length ~ Sepal.Width, data = iris)
#>
#> Residuals:
#>     Min      1Q  Median      3Q     Max
#> -1.5561 -0.6333 -0.1120  0.5579  2.2226
#>
#> Coefficients:
#>             Estimate Std. Error t value Pr(>|t|)
#> (Intercept)   6.5262     0.4789   13.63   <2e-16 ***
#> Sepal.Width  -0.2234     0.1551   -1.44    0.152
#> ---
#> Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
#>
#> Residual standard error: 0.8251 on 148 degrees of freedom
#> Multiple R-squared:  0.01382,   Adjusted R-squared:  0.007159
#> F-statistic: 2.074 on 1 and 148 DF,  p-value: 0.1519
```

the same as

```
summary(lm_iris)

#>
#> Call:
#> lm(formula = Sepal.Length ~ Sepal.Width, data = iris)
#>
#> Residuals:
#>     Min      1Q  Median      3Q     Max
#> -1.5561 -0.6333 -0.1120  0.5579  2.2226
#>
#> Coefficients:
#>             Estimate Std. Error t value Pr(>|t|)
#> (Intercept)   6.5262     0.4789   13.63   <2e-16 ***
#> Sepal.Width  -0.2234     0.1551   -1.44    0.152
#> ---
#> Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
#>
#> Residual standard error: 0.8251 on 148 degrees of freedom
#> Multiple R-squared:  0.01382,   Adjusted R-squared:  0.007159
#> F-statistic: 2.074 on 1 and 148 DF,  p-value: 0.1519
```

One could coerce a different method upon a different class, for example using `summary.data.frame` on an "lm" object:

```
summary.data.frame(lm_iris)

#>   coefficients      residuals         effects              rank
#>  Min.  :-0.2234  Min.  :-1.5561  Min.  :-71.56593  Min.   :2
#>  1st Qu.: 1.4640  1st Qu.:-0.6333  1st Qu.: -0.65192  1st Qu.:2
#>  Median : 3.1514  Median :-0.1120  Median : -0.00897  Median :2
#>  Mean  : 3.1514  Mean  : 0.0000  Mean   : -0.42040  Mean   :2
#>  3rd Qu.: 4.8388  3rd Qu.: 0.5579  3rd Qu.:  0.61051  3rd Qu.:2
#>  Max.  : 6.5262  Max.  : 2.2225  Max.   :  2.15225  Max.   :2
#>  fitted.values      assign     qr.Length qr.Class qr.Mode df.residual
#>  Min.  :5.543  Min.  :0.00    300        -none-   numeric  Min.   :148
#>  1st Qu.:5.789  1st Qu.:0.25    2         -none-   numeric  1st Qu.:148
#>  Median :5.856  Median :0.50    2         -none-   numeric  Median :148
#>  Mean  :5.843  Mean  :0.50    1          -none-   numeric  Mean   :148
#>  3rd Qu.:5.901  3rd Qu.:0.75    1         -none-   numeric  3rd Qu.:148
#>  Max.  :6.080  Max.  :1.00                                 Max.   :148
#>  xlevels       call       terms
```

```
#> Length:0      Length:3     Length:3
#> Class :list   Class :call  Class1:terms
#> Mode  :list   Mode  :call  Class2:formula
#>                             Mode  :call
#>
#>
#> model.Sepal.Length  model.Sepal.Width
#> Min.   :4.300000    Min.   :2.000000
#> 1st Qu.:5.100000    1st Qu.:2.800000
#> Median :5.800000    Median :3.000000
#> Mean   :5.843333    Mean   :3.057333
#> 3rd Qu.:6.400000    3rd Qu.:3.300000
#> Max.   :7.900000    Max.   :4.400000
```

However the output may be a bit confusing.

To summarize, the important feature of S3 methods worth noting is that only the **first part**, summary, is required to be used on these objects of different classes. If you would like to find all of the different classes an S3 method operates on, use the command, methods(summary).

Why hide the text?

Hiding the trailing text after the . avoids the need to use a different summary function for every class. This means that one does not need to remember to use summary.lm for linear models, or summary.data.frame for data frames, or summary.aProposterousClassOfObject. By using S3 methods, cognitive load is reduced - you don't have to think as much to remember what class an object is - and the commands are more intuitive. To get a summary of most objects, use summary, to plot most objects, use plot. Perhaps the most nifty feature of all is that a user can create their own S3 methods using the same functions such as summary and plot. This means a user can create their own special class of object and then write their own S3 method for it - e.g., summary.myclass or plot.myclass, each proiding appropriate summary information, or nice plots, for that object.

How to make your own S3 method?

Creating your own S3 method is not particularly difficult and is often highly practical. A use case scenario for creating an S3 method is now discussed.

The Residual Sums of Squares (RSS), $\sum(Y_i - \hat{Y})^2$ is a useful metric for determining model accuracy for continuous outcomes. For example, for a Classification and Regression Tree

```
library(rpart)
```

```
fit.rpart <- rpart(Sepal.Width ~ Sepal.Length + Petal.Length + Petal.Width + Species, data = iris)
```

The RSS is calculated as

```
print_rss <- sum(residuals(fit.rpart)**2)
```

One might be inclined to write a function to perform this task

```
rss <- function(x){

  sum(residuals(x)**2)

}
```

```
rss(fit.rpart)
```

```
#> [1] 10.17245
```

However, what to do when there are many different decision tree models that one would like to compare, say boosted regression trees (BRT), and random forests (RF). The same code will not work:

```
library(randomForest)
```

```
#> randomForest 4.6-12
```

```
#> Type rfNews() to see new features/changes/bug fixes.

set.seed(71)
fit.rf <- randomForest(Sepal.Length ~ ., data=iris, importance=TRUE,
                       proximity=TRUE)

rss(fit.rf)

#> [1] 0
```

In this case, one could write three functions, one for each decision tree method: "rss_rpart", "rss_brt", and "rss_rf". But to avoid having three functions and instead use just one, one could place all three functions inside of one function, using an if-then-else clause to direct the object of the appropriate class to the appropriate method. This is what I shall call a "Poor man's S3 method".

\begin{Schunk} \begin{Sinput} dt_rss <- function (x){

if ("rpart" %in% class(x)) {

result <- sum((residuals(x)**2))

return(result)

}

else if ("gbm" %in% class(x)) {

result <- sum(x$residuals**n2)

return(result)

}

else if ("randomForest" %in% class(x)) {

temp <- $xy - x$predicted

result <- sum(temp**2)

return(result)

}

else warning(paste(class(x), "is not of an rpart, gbm, or randomForest object")) } \end{Sinput} \end{Schunk}

Here it is in action

```
dt_rss(fit.rpart)

#> [1] 10.17245
```

The RSS method works, and if it is applied to a class that is not known, a special message is provided

```
fit.lm <- lm(Sepal.Width ~ Species, data = iris)

dt_rss(fit.lm)

#> Warning in dt_rss(fit.lm): lm is not of an rpart, gbm, or randomForest
#> object
```

The "poor man's S3 method" does what it needs to do. However, the cluster of ifelses is more difficult to read, manage, and debug, and is not as elegant and not as modular as a real S3 method. So let us create an S3 method.

First the S3 method is defined using UseMethod(), which creates the building block of an S3 method, the "root", if you will.

```
rss <- function(x) UseMethod("rss")
```

Here we have specified that our method will be called rss. Now we need to create the special cases of rss - the methods rss.rpart, rss.gbm, and rss.randomForest, where the sections of code after rss. are the classes of object we want them to work on.

\begin{Schunk} \begin{Sinput} rss.rpart <- function(x){

sum((residuals(x)**2))

}

rss.gbm <- function(x){

```
    sum(x$residuals**2)
    }
    rss.randomForest <- function(x){
    res <- xy − xpredicted
    sum(res**2)
    } \end{Sinput} \end{Schunk}
```

A default method can also be created - `rss.default` - which, as the name suggests, is the default method when other classes are not present.

```
rss.default <- function(x, ...){

  warning(paste("RSS does not know how to handle object of class ", class(x), "and can only be used on classe

          }
```

In this case a warning is issued, to let the user know that the object class they were using was not appropriate.

We can now apply the `rss` method to an `rpart` model

```
rss(fit.rpart)

#> [1] 10.17245
```

Also observe what happens when the object used is not of the decision tree classes

```
rss(lm.fit)

#> Warning in rss.default(lm.fit): RSS does not know how to handle object of
#> class function and can only be used on classes rpart, gbm, and randomForest
```

This guide to S3 methods was written to provide R users with the minimal amount of information to start building their own S3 methods. For a more complete treatment on S3 methods, see Advanced-R, R Packages, and this blog, this resource.

Extras

For the uninitiated, you may find the class of an object using the command, `class()`, on the object. For example:

```
x <- c(1, 2, 3, 4, 5)

x

#> [1] 1 2 3 4 5

class(x)

#> [1] "numeric"
```

Here, showing that the object `x` is of class `numeric`.

`str()` can also provide more information:

```
str(x)

#>  num [1:5] 1 2 3 4 5
```

In this case, revealing that `x` is `numeric`, showing its contents.

Idea

Idea: Make a cheatsheet / infographic for writing functions in R, and for making s3 methods. Let's call it "The Anatomy of S3 Methods"

Nicholas Tierney
Queensland University of Technology
Level 8, Y Block, Main Drive, QUT, Brisbane, Australia
nicholas.tierney@gmail.com